

# Assessing Product Recommendation Likelihood

May 11, 2018

- 1 The dataset consists of Women's fashion online shop reviews, consisting of a title, a review text, and whether the review author would recommend the product. The goal is to determine whether a reviewer will recommend a product or not based on review title and review. Such a classifier can help find out what is good or bad about certain products or to highlight more relevant reviews (like a very critical and a very positive one) about certain products.

```
In [1]: import numpy as np
import scipy.io as sc
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.pipeline import make_pipeline
import matplotlib.pyplot as plt
from sklearn.preprocessing import Normalizer
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
%matplotlib inline
```

## 1.1 1. Data loading and pre-processing

```
In [2]: trainDat = pd.read_csv("data_train.csv")
testDat = pd.read_csv("data_test.csv")
```

```
In [3]: print(trainDat.shape)
print(testDat.shape)
```

(17614, 3)

(5872, 3)

```
In [4]: trainDat.head(n=10)
```

```
Out[4]:
```

	Title \
0	Beautiful unique dress

```

1         Had high hopes but...
2     Buttons -buttons so cute!
3         Love this dress
4     Perfect summer pants!
5         Too bulky
6         Love it
7 Beautiful drape, very deep v
8         NaN
9     Awkward length in front

```

	Review	Recommended
0	Wore this to my sons wedding. found it last mi...	1
1	Gals, if you absolutely must have this top wai...	0
2	I hardly believe i have not reviewed htis yet...	1
3	This dress is so cute and comfortable. i'm nor...	1
4	The linen- cotton blend breathes so well for a...	1
5	I like the quality of material of this top but...	0
6	Great top. i got the size 4 but i ordered a 6 ...	1
7	I bought this yesterday and love it. it drapes...	1
8	The design of the pant hangs nicely and fit,well,	1
9	It seems like most people love this tee, so yo...	0

```
In [5]: trainDat.columns
```

```
Out[5]: Index(['Title', 'Review', 'Recommended'], dtype='object')
```

```
In [6]: #Dropping any rows with missing information (with respect to any of the three columns)
trainDatNew = trainDat.dropna(axis=0,how='any')
trainDatNew.shape
```

```
Out[6]: (14762, 3)
```

```
In [7]: text_train_title = trainDatNew['Title'].tolist()
text_train_review = trainDatNew['Review'].tolist()
```

```
In [8]: text_train_title[:5]
```

```
Out[8]: ['Beautiful unique dress',
        'Had high hopes but...',
        'Buttons -buttons so cute!',
        'Love this dress',
        'Perfect summer pants!']
```

```
In [9]: text_train_review[:5]
```

```
Out[9]: ['Wore this to my sons wedding. found it last minute when i had changed my mind about a
        "Gals, if you absolutely must have this top wait for it to go on sale (don't worry, i
        'I hardly believe i have not reviewed htis yet... so i tried it on in the sotre and w
        "This dress is so cute and comfortable. i'm normally a 0 or 2 and the xs is plenty ro
        "The linen- cotton blend breathes so well for a hot summer day! the design on the wai
```

```

In [10]: typesList = []
         for a,b in zip(text_train_title,text_train_review):
             typePair = (type(a),type(b))
             typesList.append(typePair)
         typesList[:5]

Out[10]: [(str, str), (str, str), (str, str), (str, str), (str, str)]

In [11]: text_train_title_review = [a+" "+b for a,b in
                                     zip(text_train_title,text_train_review)]
         text_train_title_review[:5]

Out[11]: ['Beautiful unique dress Wore this to my sons wedding. found it last minute when i had
          "Had high hopes but... Gals, if you absolutely must have this top wait for it to go on
          'Buttons -buttons so cute! I hardly believe i have not reviewed htis yet... so i tried
          "Love this dress This dress is so cute and comfortable. i'm normally a 0 or 2 and the
          "Perfect summer pants! The linen- cotton blend breathes so well for a hot summer day

In [12]: y_train = trainDatNew['Recommended']
         trainDatNew = trainDatNew.drop(columns=['Recommended'])

In [13]: #check if any value missing in training labels
         print(y_train.isnull().values.any())
         print(len(y_train))

False
14762

In [14]: testDatNew = testDat.dropna(axis=0,how='any')
         testDatNew.shape

Out[14]: (4913, 3)

In [15]: testDatNew.columns

Out[15]: Index(['Title', 'Review', 'Recommended'], dtype='object')

In [16]: y_test = testDatNew['Recommended']
         testDatNew = testDatNew.drop(columns=['Recommended'])
         text_test_title = testDatNew['Title'].tolist()
         text_test_review = testDatNew['Review'].tolist()

In [17]: print(len(y_test))

4913

In [18]: text_test_title_review = [a+" "+b for a,b in
                                     zip(text_test_title,text_test_review)]
         text_test_title_review[:5]

```

```
Out[18]: ['Perfect pair!! This pair of age stevie capris is everything that i could want in a p
"Talk about creature comforts! This is a beautifully designed jacket that everyone w
'So comfotable I ordered this dress in Op since i am 5ft. it fits great its not too s
'Love this cozy sweater Cozy and nice quality. great fit and runs true to size. coul
"Don't miss this top This is a great top and worth the money. nice flow to it, soft m
```

## 1.2 2. Basic Model Building and Evaluation

### 1.2.1 Model building without GridSearchCV

#### Evaluation Using Title field only

```
In [19]: vect1 = CountVectorizer()
X_train_title = vect1.fit_transform(text_train_title)
X_val_title = vect1.transform(text_test_title)
print(len(vect1.get_feature_names()))
```

3219

```
In [20]: X_train_title
```

```
Out[20]: <14762x3219 sparse matrix of type '<class 'numpy.int64'>'
with 47745 stored elements in Compressed Sparse Row format>
```

```
In [21]: feature_names_title = vect1.get_feature_names()
print(feature_names_title[:10])
print(feature_names_title[1000:1010])
print(feature_names_title[:500])
```

```
['00', '10', '100', '12', '14', '150', '16', '16s', '1950', '1960s']
['farm', 'farrah', 'fashion', 'fashionable', 'fast', 'fat', 'fatal', 'fav', 'fave', 'favorite']
['00', 'chests', 'farm', 'kick', 'petites', 'slit', 'use']
```

```
In [22]: from sklearn.linear_model import LogisticRegressionCV
lr1 = LogisticRegressionCV(scoring='roc_auc').fit(X_train_title, y_train)
```

```
In [23]: lr1.C_
```

```
Out[23]: array([2.7825594])
```

```
In [24]: lr1.score(X_val_title, y_test)
```

```
Out[24]: 0.8827600244249949
```

```
In [25]: def plot_important_features(coef, feature_names, top_n=20,
ax=None, rotation=60):
    if ax is None:
        ax = plt.gca()
    inds = np.argsort(coef)
```

```

low = inds[:top_n]
high = inds[-top_n:]
important = np.hstack([low, high])
myrange = range(len(important))
colors = ['red'] * top_n + ['blue'] * top_n

ax.bar(myrange, coef[important], color=colors)
ax.set_xticks(myrange)
ax.set_xticklabels(feature_names[important],
                    rotation=rotation, ha="right")
ax.set_xlim(-.7, 2 * top_n)
ax.set_frame_on(False)

```

```

In [26]: %%time
from sklearn.model_selection import cross_val_score
bow_pipe_title = make_pipeline(CountVectorizer(),
                               LogisticRegressionCV(scoring='roc_auc'))
print(cross_val_score(bow_pipe_title, text_train_title,
                      y_train, cv=5, scoring='roc_auc'))

```

```
[0.91901162 0.9244011 0.91295734 0.92318662 0.92155286]
```

```
CPU times: user 12.2 s, sys: 119 ms, total: 12.3 s
```

```
Wall time: 12.8 s
```

```
In [27]: bow_pipe_title.fit(text_train_title,y_train)
```

```

lr_title = bow_pipe_title.named_steps['logisticregressioncv']
feature_names_title = np.array(bow_pipe_title
                               .named_steps['countvectorizer']
                               .get_feature_names())

#n_classes_title = len(lr_title.classes_)

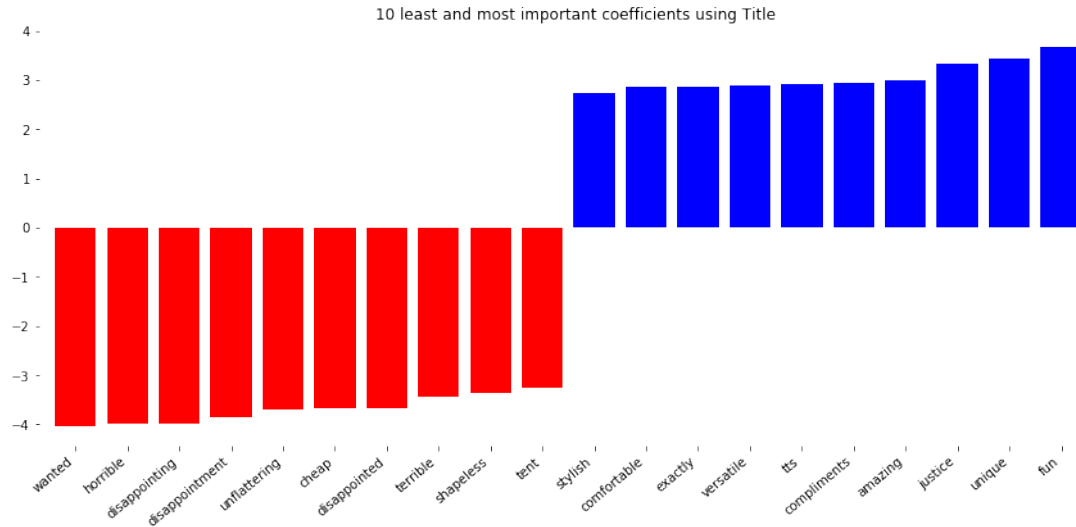
```

```

In [28]: plt.figure(figsize=(15, 6))
plot_important_features(lr_title.coef_.ravel(),
                        np.array(feature_names_title),
                        top_n=10, rotation=40)

ax = plt.gca()
plt.title("10 least and most important coefficients using Title")
plt.show()

```



```
In [29]: cv_title = bow_pipe_title.named_steps['countvectorizer']
X_train_title = cv_title.transform(text_train_title)
X_val_title = cv_title.transform(text_test_title)
lr_title.score(X_val_title, y_test)
```

```
Out[29]: 0.8827600244249949
```

```
In [30]: print(X_train_title.shape)
print(len(feature_names_title))
```

```
(14762, 3219)
```

```
3219
```

### Evaluation using Review field only

```
In [31]: %%time
bow_pipe_review = make_pipeline(CountVectorizer(),
                                LogisticRegressionCV(scoring='roc_auc'))
print(cross_val_score(bow_pipe_review, text_train_review,
                      y_train, cv=5, scoring='roc_auc'))
```

```
[0.92095963 0.9270255 0.9173906 0.92501866 0.92596217]
```

```
CPU times: user 1min 11s, sys: 1.33 s, total: 1min 12s
```

```
Wall time: 39.9 s
```

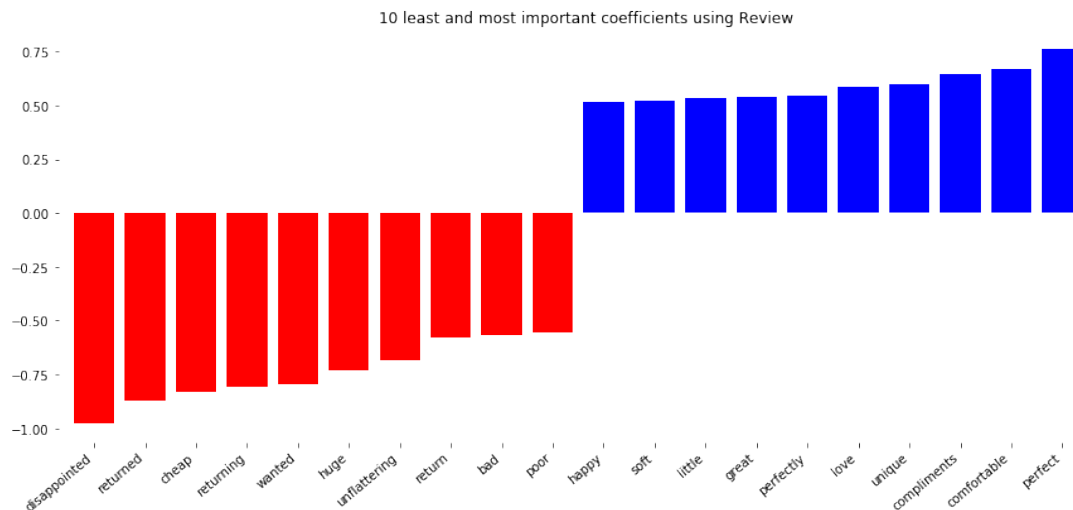
```
In [32]: bow_pipe_review.fit(text_train_review,y_train)
```

```
lr_review = bow_pipe_review.named_steps['logisticregressioncv']
```

```
feature_names_review = np.array(bow_pipe_review
                                .named_steps['countvectorizer']
                                .get_feature_names())
```

```
In [33]: plt.figure(figsize=(15, 6))
plot_important_features(lr_review.coef_.ravel(),
                        np.array(feature_names_review),
                        top_n=10, rotation=40)

ax = plt.gca()
plt.title("10 least and most important coefficients using Review")
plt.show()
```



```
In [34]: cv_review = bow_pipe_review.named_steps['countvectorizer']
X_train_review = cv_review.transform(text_train_review)
X_val_review = cv_review.transform(text_test_review)
lr_review.score(X_val_review, y_test)
```

```
Out[34]: 0.8843883574190922
```

```
In [35]: print(X_train_review.shape)
print(len(feature_names_review))
```

```
(14762, 12041)
12041
```

## Evaluation using Title and Review fields

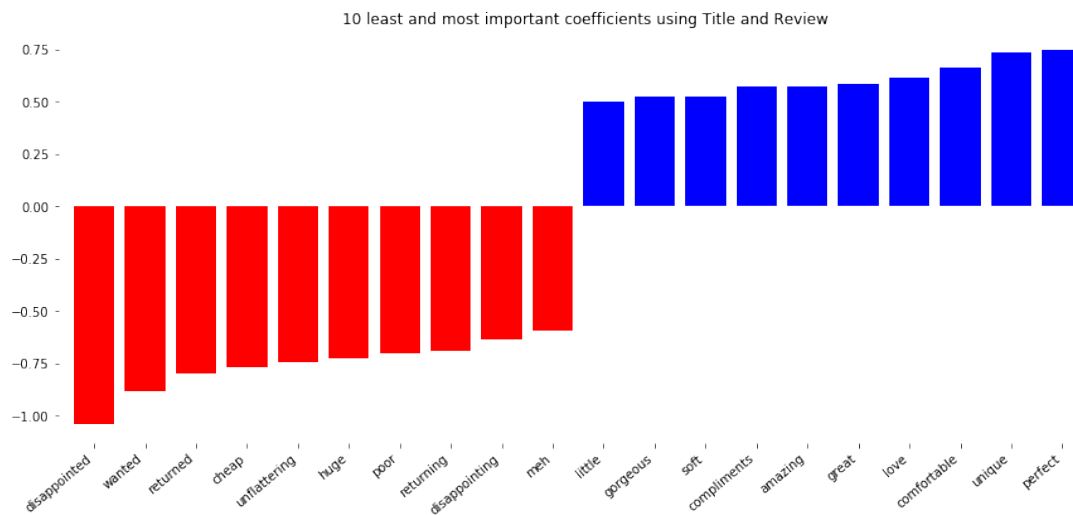
```
In [36]: %%time
bow_pipe_title_review = make_pipeline(CountVectorizer(),
                                      LogisticRegressionCV(scoring='roc_auc'))
cross_val_score(bow_pipe_title_review, text_train_title_review,
                y_train, cv=5, scoring='roc_auc')
```

CPU times: user 1min 19s, sys: 1.82 s, total: 1min 21s  
Wall time: 46.4 s

```
In [37]: bow_pipe_title_review.fit(text_train_title_review,y_train)
```

```
lr_title_review = bow_pipe_title_review.named_steps['logisticregressioncv']  
feature_names_title_review = np.array(bow_pipe_title_review  
                                       .named_steps['countvectorizer']  
                                       .get_feature_names())
```

```
In [38]: plt.figure(figsize=(15, 6))  
plot_important_features(lr_title_review.coef_.ravel(),  
                       np.array(feature_names_title_review),  
                       top_n=10, rotation=40)  
  
ax = plt.gca()  
plt.title("10 least and most important coefficients using Title and Review")  
plt.show()
```



```
In [39]: cv_title_review = bow_pipe_title_review.named_steps['countvectorizer']  
X_train_title_review = cv_title_review.transform(text_train_title_review)  
X_val_title_review = cv_title_review.transform(text_test_title_review)  
lr_title_review.score(X_val_title_review, y_test)
```

```
Out[39]: 0.8963973132505597
```

```
In [40]: print(X_train_title_review.shape)  
         print(len(feature_names_title_review))
```

```
(14762, 12515)  
12515
```



## Vectorizing Title and Review individually and concatenating the vector representations

```
In [41]: from scipy.sparse import hstack
```

```
feature_names_title_new = [feat+"_t" for feat in feature_names_title]
print(feature_names_title_new[:10])
```

```
feature_names_review_new = [feat+"_r" for feat in feature_names_review]
print(feature_names_review_new[:10])
```

```
['00_t', '10_t', '100_t', '12_t', '14_t', '150_t', '16_t', '16s_t', '1950_t', '1960s_t']
['00_r', '00p_r', '03_r', '03dd_r', '04_r', '06_r', '0dd_r', '0in_r', '0p_r', '0petite_r']
```

```
In [42]: print(X_train_title.shape)
print(X_train_review.shape)
X_train_title_review_concat = hstack([X_train_title, X_train_review])
print(X_train_title_review_concat.shape)
```

```
(14762, 3219)
(14762, 12041)
(14762, 15260)
```

```
In [43]: %%time
print(cross_val_score(LogisticRegressionCV(), X_train_title_review_concat,
                      y_train, cv=5, scoring='roc_auc'))
```

```
[0.9454991  0.94365854 0.93439082 0.94418817 0.94631581]
CPU times: user 1min 15s, sys: 1.36 s, total: 1min 16s
Wall time: 39.3 s
```

```
In [45]: feature_names_title_review_concat = feature_names_title_new + \
        feature_names_review_new
print(len(feature_names_title_review_concat))
```


```
15260
```

```
In [46]: lr_title_review_concat = LogisticRegressionCV(scoring='roc_auc')
lr_title_review_concat.fit(X_train_title_review_concat, y_train)
```

```
Out[46]: LogisticRegressionCV(Cs=10, class_weight=None, cv=None, dual=False,
    fit_intercept=True, intercept_scaling=1.0, max_iter=100,
    multi_class='ovr', n_jobs=1, penalty='l2', random_state=None,
    refit=True, scoring='roc_auc', solver='lbfgs', tol=0.0001,
    verbose=0)
```

```
In [47]: plt.figure(figsize=(15, 6))
          plot_important_features(lr_title_review_concat.coef_.ravel(),
                                np.array(feature_names_title_review_concat),
                                top_n=10, rotation=40)

          ax = plt.gca()
          plt.title("10 least and most important coefficients \
                    using Title and Review Concatenated")
          plt.show()
```



```
In [48]: print(X_val_title.shape)
          print(X_val_review.shape)
          X_val_title_review_concat = hstack([X_val_title,X_val_review])
          print(X_val_title_review_concat.shape)
```

(4913, 3219)

(4913, 12041)

(4913, 15260)

```
In [49]: lr_title_review_concat.score(X_val_title_review_concat,y_test)
```

Out[49]: 0.9037248117239975

## 1.2.2 With GridSearchCV

### Evaluating using Title field only

```
In [50]: from sklearn.linear_model import LogisticRegression
```

```
In [53]: %%time
param_grid_lr = {"logisticregression__C": [100, 10, 1, 0.1, 0.001]}
grid_title = GridSearchCV(make_pipeline(CountVectorizer(),
                                      LogisticRegression()),
                          param_grid=param_grid_lr, cv=5,
                          scoring="roc_auc")
grid_title.fit(text_train_title,y_train)

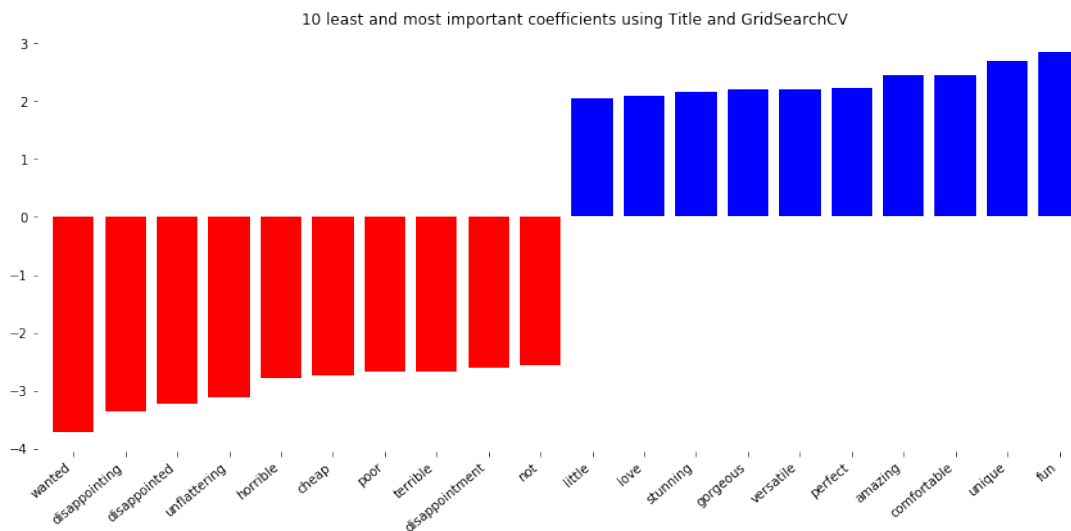
print(grid_title.best_score_)
print(grid_title.best_params_)

lr_title_grid = grid_title.best_estimator_.named_steps['logisticregression']
feature_names_title_grid = np.array(grid_title
                                    .best_estimator_
                                    .named_steps['countvectorizer']
                                    .get_feature_names())

0.9204469952584714
{'logisticregression__C': 1}
CPU times: user 5.12 s, sys: 60.5 ms, total: 5.18 s
Wall time: 5.22 s
```

```
In [54]: plt.figure(figsize=(15, 6))
plot_important_features(lr_title_grid.coef_.ravel(),
                      np.array(feature_names_title_grid),
                      top_n=10, rotation=40)

ax = plt.gca()
plt.title("10 least and most important coefficients \
using Title and GridSearchCV")
plt.show()
```



```
In [55]: print(grid_title.score(text_test_title,y_test))
```

0.9214984071094153

### Evaluating using Review field only

```
In [57]: %%time
grid_review = GridSearchCV(make_pipeline(CountVectorizer(),
                                         LogisticRegression()),
                           param_grid=param_grid_lr, cv=5,
                           scoring="roc_auc")
grid_review.fit(text_train_review,y_train)

print(grid_review.best_score_)
print(grid_review.best_params_)

lr_review_grid = grid_review.best_estimator_.named_steps['logisticregression']
feature_names_review_grid = np.array(grid_review
                                     .best_estimator_
                                     .named_steps['countvectorizer']
                                     .get_feature_names())
```

0.9237641674362176

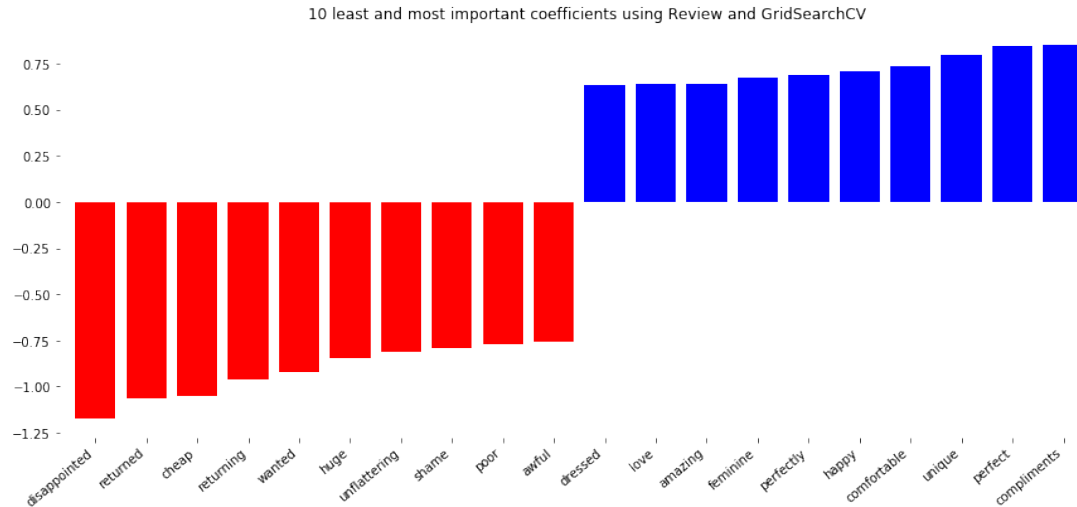
{'logisticregression\_\_C': 0.1}

CPU times: user 1min 17s, sys: 1.75 s, total: 1min 19s

Wall time: 57.8 s

```
In [58]: plt.figure(figsize=(15, 6))
plot_important_features(lr_review_grid.coef_.ravel(),
                        np.array(feature_names_review_grid),
                        top_n=10, rotation=40)

ax = plt.gca()
plt.title("10 least and most important coefficients using \
Review and GridSearchCV")
plt.show()
```



```
In [59]: print(grid_review.score(text_test_review,y_test))
```

0.9221417324054474

### Evaluating using Title and Review

```
In [60]: %%time
```

```
grid_title_review = GridSearchCV(make_pipeline(CountVectorizer(),
                                                LogisticRegression()),
                                param_grid=param_grid_lr, cv=5,
                                scoring="roc_auc")
grid_title_review.fit(text_train_title_review,y_train)
```

```
print(grid_title_review.best_score_)
print(grid_title_review.best_params_)
```

```
lr_title_review_grid = grid_title_review.best_estimator_.named_steps['logisticregressor']
feature_names_title_review_grid = np.array(grid_title_review
                                             .best_estimator_
                                             .named_steps['countvectorizer']
                                             .get_feature_names())
```

0.9424097253719665

```
{'logisticregression__C': 0.1}
```

CPU times: user 1min 14s, sys: 1.69 s, total: 1min 16s

Wall time: 56.7 s

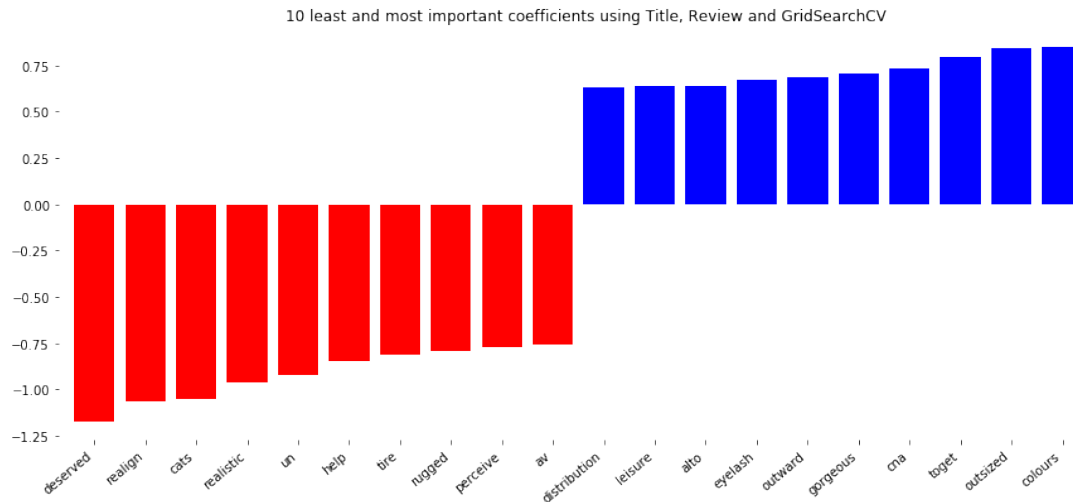
```
In [61]: plt.figure(figsize=(15, 6))
         plot_important_features(lr_review_grid.coef_.ravel(),
```

```

np.array(feature_names_title_review_grid),
top_n=10, rotation=40)

ax = plt.gca()
plt.title("10 least and most important coefficients \
using Title, Review and GridSearchCV")
plt.show()

```



```

In [62]: print(grid_title_review.score(text_test_title_review,y_test))

0.9388955774186873

```

### Vectorizing Title and Review individually and concatenating the vector representations

```

In [63]: cv_title_grid = grid_title.best_estimator_.named_steps['countvectorizer']
cv_review_grid = grid_review.best_estimator_.named_steps['countvectorizer']
X_train_title_grid = cv_title_grid.transform(text_train_title)
X_train_review_grid = cv_review_grid.transform(text_train_review)
X_train_title_review_concat_grid = hstack([X_train_title_grid,
                                           X_train_review_grid])

X_val_title_grid = cv_title_grid.transform(text_test_title)
X_val_review_grid = cv_review_grid.transform(text_test_review)
X_val_title_review_concat_grid = hstack([X_val_title_grid,
                                          X_val_review_grid])

In [67]: feature_names_title_grid_new = [feat+"_t" for
                                          feat in feature_names_title_grid]
feature_names_review_grid_new = [feat+"_r" for
                                  feat in feature_names_review_grid]
feature_names_title_review_concat_grid = feature_names_title_grid_new + \
                                          feature_names_review_grid_new

```

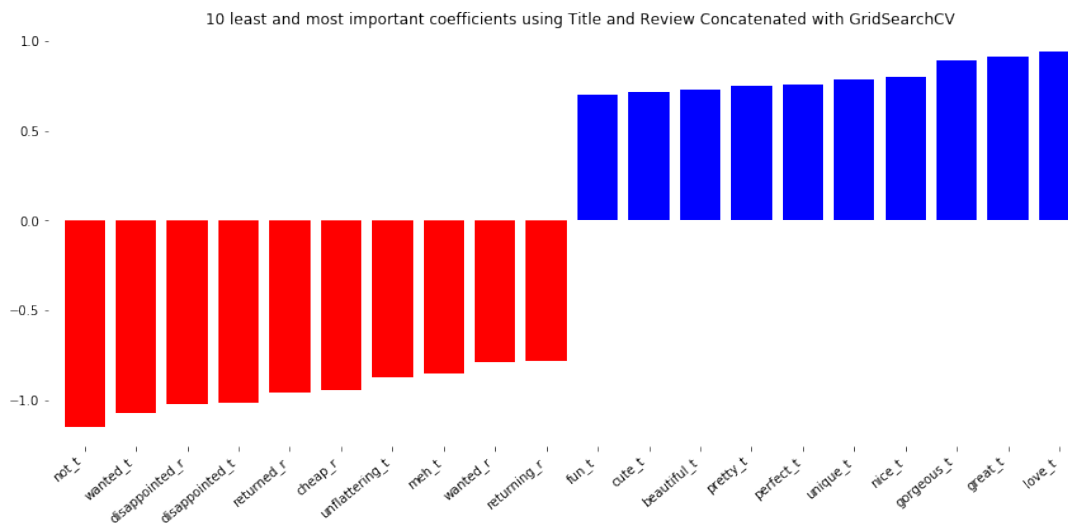
```
In [68]: %%time
param_grid_lr_new = {"C": [100, 10, 1, 0.1, 0.001]}
grid_title_review_concat = GridSearchCV(LogisticRegression(),
                                         param_grid=param_grid_lr_new,
                                         cv=5, scoring="roc_auc")
grid_title_review_concat.fit(X_train_title_review_concat_grid,y_train)

lr_title_review_grid_concat = grid_title_review_concat.best_estimator_
print(grid_title_review_concat.best_score_)
print(grid_title_review_concat.best_params_)

0.9436040159028604
{'C': 0.1}
CPU times: user 35 s, sys: 966 ms, total: 35.9 s
Wall time: 20.7 s
```

```
In [69]: plt.figure(figsize=(15, 6))
plot_important_features(lr_title_review_grid_concat.coef_.ravel(),
                       np.array(feature_names_title_review_concat_grid),
                       top_n=10, rotation=40)

ax = plt.gca()
plt.title("10 least and most important coefficients using \
Title and Review Concatenated with GridSearchCV")
plt.show()
```



```
In [70]: print(grid_title_review_concat.score(X_val_title_review_concat_grid,
                                              y_test))
```

0.9431201401345355

The 4th setting which is vectorizing title and reviews individually and concatenating them works the best marginally. The best regularization parameter (C) for this setting is 0.1.

### 1.3 3. Feature Tuning

Using TfidfVectorizer to check its impact on model score and importance coefficients

```
In [72]: %%time
param_grid_lr = {"logisticregression__C": [100, 10, 1, 0.1, 0.001]}
grid_tfidf = GridSearchCV(make_pipeline(TfidfVectorizer(),
                                         LogisticRegression()),
                          param_grid=param_grid_lr, cv=5,
                          scoring="roc_auc")
grid_tfidf.fit(text_train_title_review,y_train)

print(grid_tfidf.best_score_)
print(grid_tfidf.best_params_)


lr_tfidf = grid_tfidf.best_estimator_.named_steps['logisticregression']
feature_names_tfidf = np.array(grid_tfidf.best_estimator_
                               .named_steps['tfidfvectorizer']
                               .get_feature_names())

0.9498551621199149
{'logisticregression__C': 1}
CPU times: user 48 s, sys: 1.23 s, total: 49.2 s
Wall time: 42.5 s

In [73]: plt.figure(figsize=(15, 6))
plot_important_features(lr_tfidf.coef_.ravel(),
                       np.array(feature_names_tfidf),
                       top_n=10, rotation=40)

ax = plt.gca()
plt.title("Most important coefficients using TfidfVectorizer")
plt.show()
```





```
In [74]: print(grid_tfidf.score(text_test_title_review,y_test))
0.9454635538445706
```

Using the TfidfVectorizer increases the test score to 0.945. Some of the important coefficients remain the same but the majority are changed.

### Using a CountVectorizer with Normalizer to check impact on model performance

```
In [77]: grid_countvect_normalized = GridSearchCV(make_pipeline(CountVectorizer(),
                                                                Normalizer(),
                                                                LogisticRegression()),
                                                param_grid=param_grid_lr, cv=5,
                                                scoring="roc_auc")

grid_countvect_normalized.fit(text_train_title_review,
                              y_train)

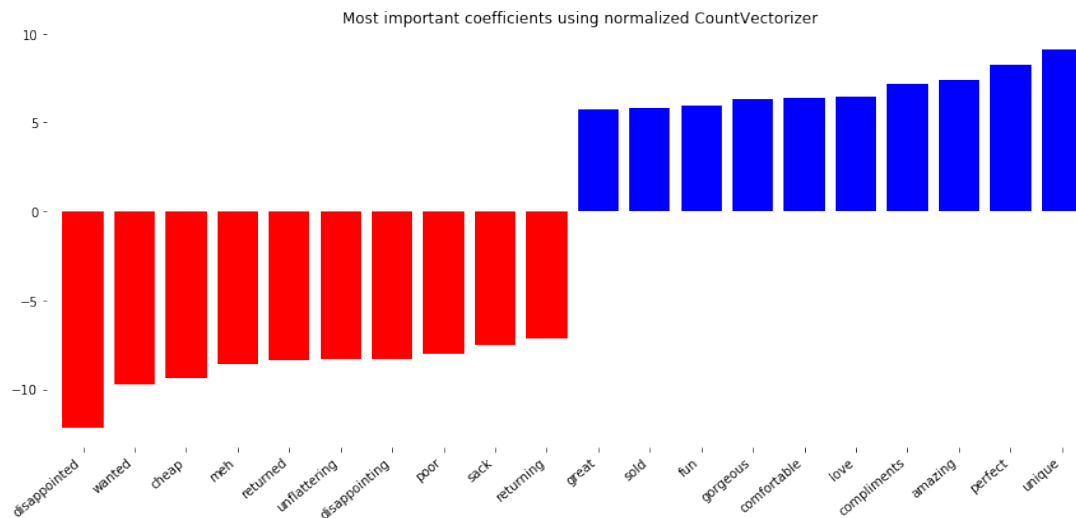
print(grid_countvect_normalized.best_score_)
print(grid_countvect_normalized.best_params_)

lr_countvect_normalized = grid_countvect_normalized\
    .best_estimator_\
    .named_steps['logisticregression']
feature_names_countvect_normalized = np.array(grid_countvect_normalized
                                                .best_estimator_\
                                                .named_steps['countvectorizer']
                                                .get_feature_names())

0.9476031785092168
{'logisticregression__C': 10}
```

```
In [78]: plt.figure(figsize=(15, 6))
        plot_important_features(lr_countvect_normalized.coef_.ravel(),
                               np.array(feature_names_countvect_normalized),
                               top_n=10, rotation=40)

        ax = plt.gca()
        plt.title("Most important coefficients using \
normalized CountVectorizer")
        plt.show()
```



```
In [79]: print(grid_countvect_normalized.score(text_test_title_review,
                                              y_test))
```

0.9441320876438387

Using a Normalizer with CountVectorizer gives a worse performance than using a TfidfVectorizer but the performance improves compared to just using the CountVectorizer without the Normalizer.

### Removing Stop Words and checking impact on number of features and model performance

```
In [81]: %%time
        grid_tfidf_stop = GridSearchCV(make_pipeline(TfidfVectorizer(stop_words =
                                                         'english'),
                                                         LogisticRegression()),
                                       param_grid=param_grid_lr, cv=5,
                                       scoring="roc_auc")
        grid_tfidf_stop.fit(text_train_title_review,y_train)

        print(grid_tfidf_stop.best_score_)
```

```

print(grid_tfidf_stop.best_params_)

lr_tfidf_stop = grid_tfidf_stop.best_estimator_.named_steps['logisticregression']
feature_names_tfidf_stop = np.array(grid_tfidf_stop
                                     .best_estimator_
                                     .named_steps['tfidfvectorizer']
                                     .get_feature_names())

0.9434995623247132
{'logisticregression__C': 1}
CPU times: user 40.4 s, sys: 805 ms, total: 41.2 s
Wall time: 35.9 s

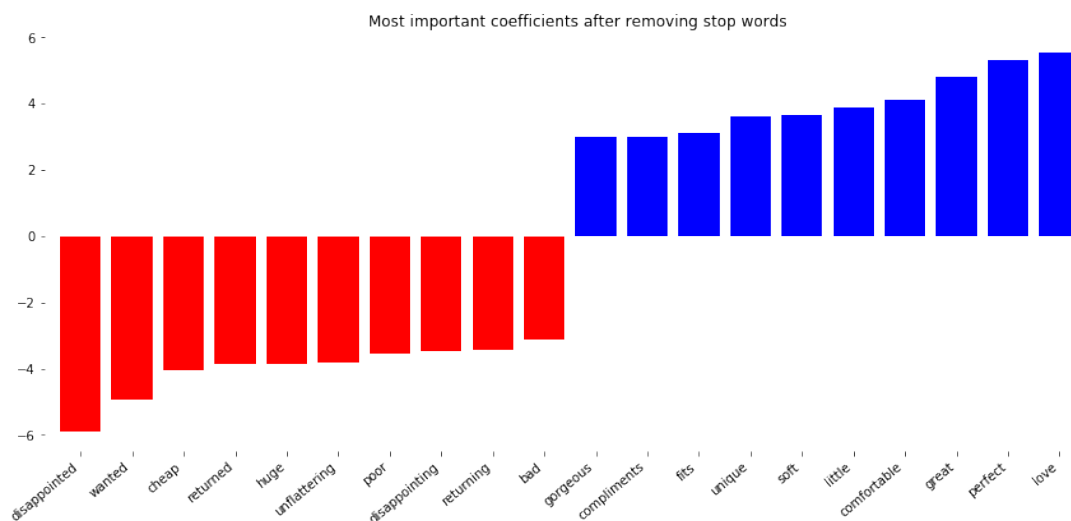
```

```

In [82]: plt.figure(figsize=(15, 6))
         plot_important_features(lr_tfidf_stop.coef_.ravel(),
                                np.array(feature_names_tfidf_stop),
                                top_n=10, rotation=40)

ax = plt.gca()
plt.title("Most important coefficients after removing stop words")
plt.show()

```



```

In [83]: print(grid_tfidf_stop.score(text_test_title_review,y_test))

0.9397407112735047

```

Standard english words as stop words don't help with the performance of the model. This is because stop words, by definition, are the most common words in english language and they will very likely be present across most documents. Since TfidfVectorizer is being used, the stop words will have less importance and hence removing them shouldn't affect the performance of the model too much. Another reason is that the total number of stop words are insignificant compared to total number of tokens.

## Limiting Vocabulary using min\_df or max\_df

```
In [85]: %%time
param_grid_lr_vocab = {"logisticregression__C": [100, 10, 1, 0.1, 0.001],
                       "tfidfvectorizer__min_df": [0.00001, 0.0001, 0.001],
                       "tfidfvectorizer__max_df": [0.5, 0.6, 0.7],
                       }

#using pipeline memory option to decrease grid-search run-time
grid_tfidf_vocab = GridSearchCV(make_pipeline(TfidfVectorizer(),
                                              LogisticRegression(),
                                              memory = 'cache_folder'),
                               param_grid=param_grid_lr_vocab, cv=5,
                               scoring="roc_auc", n_jobs = -1)
grid_tfidf_vocab.fit(text_train_title_review,y_train)

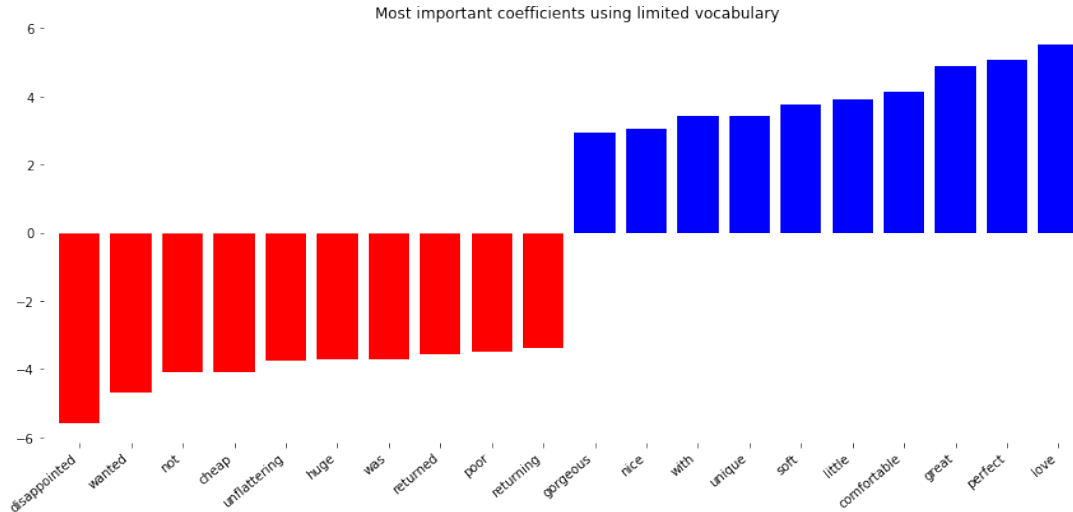
print(grid_tfidf_vocab.best_score_)
print(grid_tfidf_vocab.best_params_)

lr_tfidf_vocab = grid_tfidf_vocab.best_estimator_.named_steps['logisticregression']
feature_names_tfidf_vocab = np.array(grid_tfidf_vocab
                                     .best_estimator_
                                     .named_steps['tfidfvectorizer']
                                     .get_feature_names())

0.9503769994060268
{'logisticregression__C': 1, 'tfidfvectorizer__max_df': 0.6, 'tfidfvectorizer__min_df': 0.0001}
CPU times: user 4.26 s, sys: 2.15 s, total: 6.41 s
Wall time: 2min 45s

In [86]: plt.figure(figsize=(15, 6))
plot_important_features(lr_tfidf_vocab.coef_.ravel(),
                       np.array(feature_names_tfidf_vocab),
                       top_n=10, rotation=40)

ax = plt.gca()
plt.title("Most important coefficients using limited vocabulary")
plt.show()
```



```
In [87]: print(grid_tfidf_vocab.score(text_test_title_review,y_test))
```

```
0.9460682879220274
```

```
In [88]: print(len(feature_names_tfidf), len(feature_names_tfidf_vocab))
```

```
12515 7176
```

The number of features are decreased from 12515 to 7176 after using `min_df = 0.0001` and `max_df = 0.6` with the `tfidf` vectorizer. The score was improved by a significant amount after limiting vocabulary.

## 1.4 4. Using n-grams and assesing impact

Using current best model, changing from unigrams to n-grams of varying length

```
In [89]: %%time
param_grid_lr_ngrams = {"tfidfvectorizer__ngram_range": [(1,2),
                                                         (1,3),
                                                         (2,3)]}

grid_tfidf_ngrams = GridSearchCV(make_pipeline(TfidfVectorizer(max_df = 0.6,
                                                             min_df = 0.0001),
                                                         LogisticRegression(C = 1),
                                                         memory = 'cache_folder'),
                                param_grid=param_grid_lr_ngrams, cv=5,
                                scoring="roc_auc", n_jobs = -1)
grid_tfidf_ngrams.fit(text_train_title_review,y_train)
```

```

print(grid_tfidf_ngrams.best_score_)
print(grid_tfidf_ngrams.best_params_)

lr_tfidf_ngrams = grid_tfidf_ngrams.best_estimator_.\
named_steps['logisticregression']
feature_names_tfidf_ngrams = np.array(grid_tfidf_ngrams.
                                       best_estimator_.
                                       named_steps['tfidfvectorizer'].
                                       get_feature_names())

```

```

0.9534868951199252
{'tfidfvectorizer__ngram_range': (1, 2)}
CPU times: user 6.76 s, sys: 510 ms, total: 7.27 s
Wall time: 1min 35s

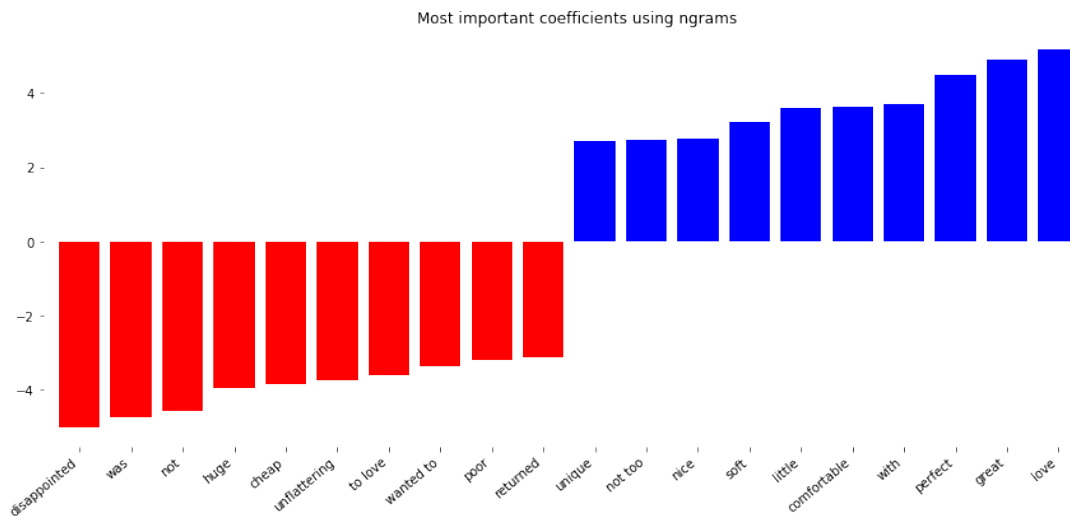
```

```

In [90]: plt.figure(figsize=(15, 6))
         plot_important_features(lr_tfidf_ngrams.coef_.ravel(),
                                np.array(feature_names_tfidf_ngrams),
                                top_n=10, rotation=40)

         ax = plt.gca()
         plt.title("Most important coefficients using ngrams")
         plt.show()

```



```

In [91]: features = pd.Series(feature_names_tfidf_ngrams)

         plt.figure(figsize=(15, 6))
         plot_important_features(np.array(pd.Series(lr_tfidf_ngrams
                                                    .coef_.ravel())[features\

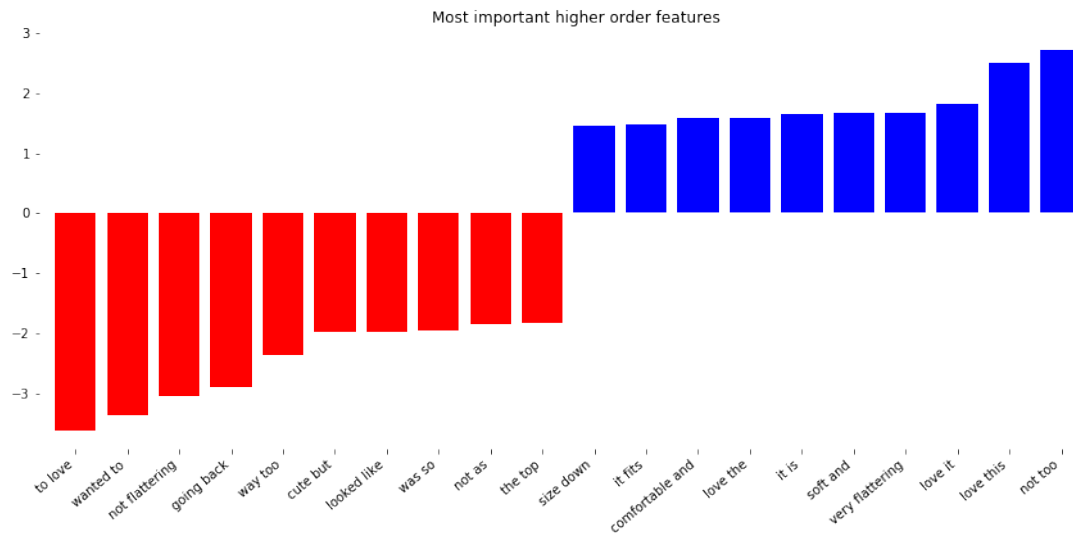
```

```

        .str
        .contains(' ')]),
    np.array(features[features.str.contains(' ')]),
    top_n=10, rotation=40)

ax = plt.gca()
plt.title("Most important higher order features")
plt.show()

```



```
In [92]: print(grid_tfidf_ngrams.score(text_test_title_review,y_test))
```

0.9518813149895486

The best performance is given my ngram\_range = (1,2) with a test score of 0.952.

### Using character n-grams and assesing impact

```
In [93]: %%time
```

```

param_grid_lr_ngrams = {"tfidfvectorizer__ngram_range": [(2,5),
                                                            (3,4),
                                                            (3,5)]}

```

```

grid_tfidf_ngrams_char = GridSearchCV(make_pipeline(TfidfVectorizer(max_df = 0.6,
                                                                    min_df = 0.0001,
                                                                    analyzer = 'char_v
                                                                    LogisticRegression(C = 1),
                                                                    memory = 'cache_folder'),
param_grid=param_grid_lr_ngrams,

```

```

cv=5,scoring="roc_auc", n_jobs = -1)

grid_tfidf_ngrams_char.fit(text_train_title_review,y_train)

print(grid_tfidf_ngrams_char.best_score_)
print(grid_tfidf_ngrams_char.best_params_)

lr_tfidf_ngrams_char = grid_tfidf_ngrams_char.best_estimator_\
.named_steps['logisticregression']
feature_names_tfidf_ngrams_char = np.array(grid_tfidf_ngrams_char
                                             .best_estimator_
                                             .named_steps['tfidfvectorizer']
                                             .get_feature_names())

0.9495127385389969
{'tfidfvectorizer__ngram_range': (3, 5)}
CPU times: user 15 s, sys: 991 ms, total: 16 s
Wall time: 3min 9s


```

```

In [94]: plt.figure(figsize=(15, 6))
         plot_important_features(lr_tfidf_ngrams_char.coef_.ravel(),
                                np.array(feature_names_tfidf_ngrams_char),
                                top_n=10, rotation=40)

         ax = plt.gca()
         plt.title("Most important coefficients using character ngrams")
         plt.show()

```



```

In [95]: print(grid_tfidf_ngrams_char.score(text_test_title_review,y_test))

0.9449816477316109

```



Even after using character n-grams, it is observed that most important features are usually whole words and they are similar to the features in word n-gram. `char_wb` is the analyzer used so as to respect the word boundaries of the character windows.

### Investigating impact of stop words and min-df on word n-grams

```
In [98]: param_grid_lr_ngrams = {"tfidfvectorizer__min_df": [0.00001,
                                                             0.0001,
                                                             0.001,
                                                             0.1],
                                "tfidfvectorizer__stop_words": ['english',
                                                                None]}

grid_tfidf_ngrams = GridSearchCV(make_pipeline(TfidfVectorizer(max_df = 0.6,
                                                             ngram_range = (1,2)),
                                           LogisticRegression(C = 1),
                                           memory = 'cache_folder'),
                                param_grid=param_grid_lr_ngrams, cv=5,
                                scoring="roc_auc", n_jobs = -1)
grid_tfidf_ngrams.fit(text_train_title_review,y_train)

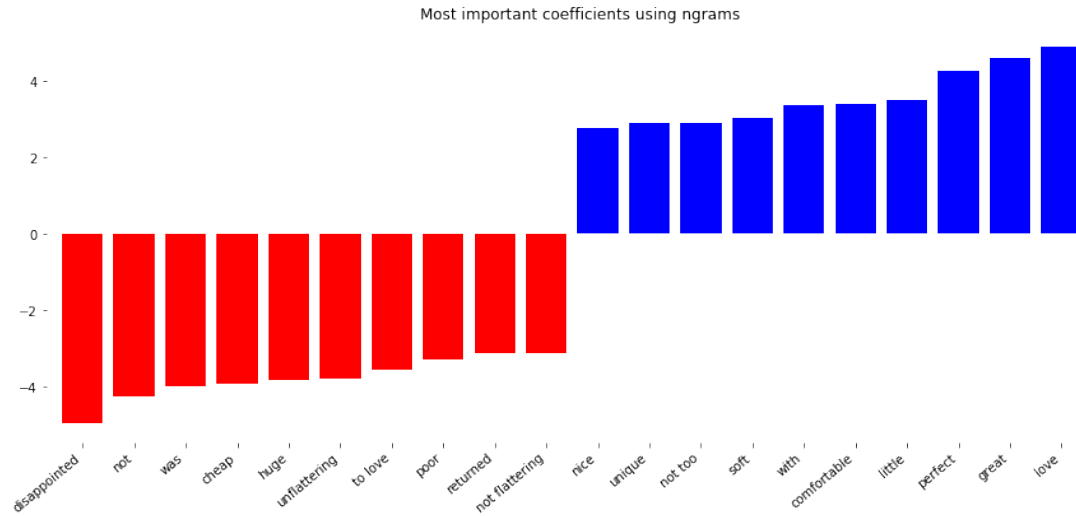
print(grid_tfidf_ngrams.best_score_)
print(grid_tfidf_ngrams.best_params_)

lr_tfidf_ngrams = grid_tfidf_ngrams.best_estimator_ \
.named_steps['logisticregression']
feature_names_tfidf_ngrams = np.array(grid_tfidf_ngrams
                                       .best_estimator_
                                       .named_steps['tfidfvectorizer']
                                       .get_feature_names())

0.9539853060627546
{'tfidfvectorizer__min_df': 0.001, 'tfidfvectorizer__stop_words': None}

In [99]: plt.figure(figsize=(15, 6))
         plot_important_features(lr_tfidf_ngrams.coef_.ravel(),
                                np.array(feature_names_tfidf_ngrams),
                                top_n=10, rotation=40)


         ax = plt.gca()
         plt.title("Most important coefficients using ngrams")
         plt.show()
```



```
In [100]: features = pd.Series(feature_names_tfidf_ngrams)

plt.figure(figsize=(15, 6))
plot_important_features(np.array(pd.Series(lr_tfidf_ngrams.coef_.ravel())[features.str.contains(' ')]),
                        np.array(features[features.str.contains(' ')]),
                        top_n=10, rotation=40)

ax = plt.gca()
plt.title("Most important higher order features")
plt.show()
```



```
In [101]: print(grid_tfidf_ngrams.score(text_test_title_review,y_test))
```

0.9523242149245934

```
In [102]: len(features)
```

```
Out[102]: 10220
```

Removing stop words gives a slightly worse performance than keeping them. And the min\_df = 0.001 gives the best score. The test score improved from 0.9518 to 0.9523. The total number of features are reduced from 12515 to 10220 after applying these transformations.

## 1.5 5. Comparing L1 and L2 penalty on best model from previous step

```
In [104]: param_grid_lr_ngrams = {"logisticregression__penalty": ['l1',
                                                                    'l2'],
                                   "logisticregression__C": [0.001,
                                                             0.01,
                                                             0.1,
                                                             1,
                                                             10,
                                                             100]}

grid_tfidf_final = GridSearchCV(make_pipeline(TfidfVectorizer(max_df = 0.6,
                                                             min_df = 0.001,
                                                             ngram_range = (1,2)),
                                       LogisticRegression(),
                                       memory = 'cache_folder'),
                               param_grid=param_grid_lr_ngrams,
                               cv=5,scoring="roc_auc", n_jobs = -1)
grid_tfidf_final.fit(text_train_title_review,y_train)

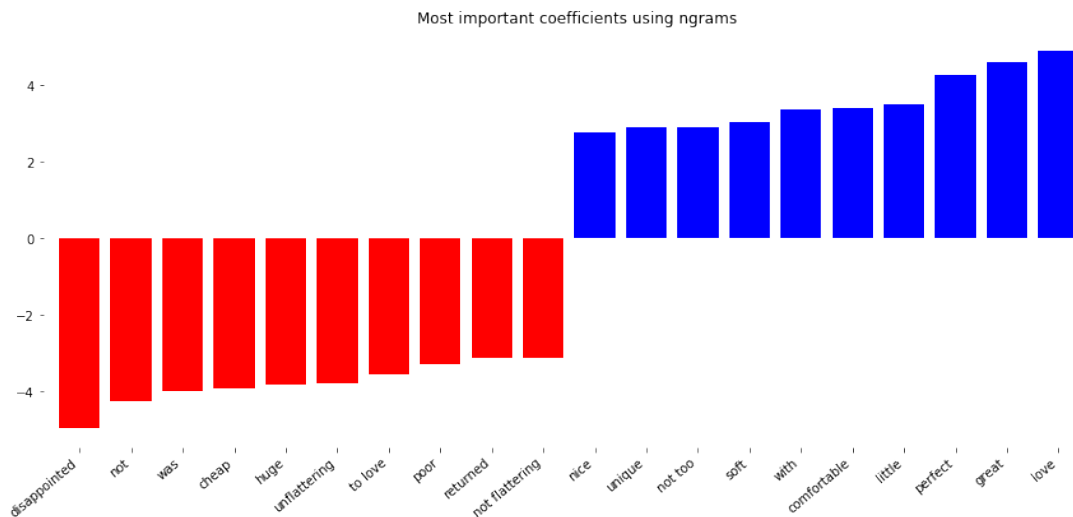
print(grid_tfidf_final.best_score_)
print(grid_tfidf_final.best_params_)

lr_tfidf_final = grid_tfidf_final.best_estimator_ \
.named_steps['logisticregression']
feature_names_tfidf_final = np.array(grid_tfidf_final
                                     .best_estimator_
                                     .named_steps['tfidfvectorizer']
                                     .get_feature_names())

0.9539853060627546
{'logisticregression__C': 1, 'logisticregression__penalty': 'l2'}

In [105]: plt.figure(figsize=(15, 6))
           plot_important_features(lr_tfidf_final.coef_.ravel(),
                                   np.array(feature_names_tfidf_final),
                                   top_n=10, rotation=40)
```

```
ax = plt.gca()
plt.title("Most important coefficients using ngrams")
plt.show()
```



```
In [106]: print(grid_tfidf_final.score(text_test_title_review,y_test))
```

```
0.9523242149245934
```

$l_2$  penalty is better than  $l_1$  penalty for this model. Different regularization parameters are tried and it is found that  $C=0.1$  is the best performing regularization parameter.

Other features that can be tried are: - Length of text - Number of *out-of-vocabulary* words - Using *sentiment* words - Presence and *frequency* of all *upper-case* words - Taking punctuation into account with character n-grams